

House Price Predictor

July 16, 2024

1 House Price Predictor using Linear Regression Model

1.1 Import All the Necessary Libraries

```
[1]: import pandas as pd import numpy as np import seaborn as
sns import matplotlib.pyplot as plt from
sklearn.linear_model import LinearRegression from
sklearn.tree import DecisionTreeRegressor from
sklearn.ensemble import RandomForestRegressor from
sklearn.preprocessing import
StandardScaler,QuantileTransformer from
sklearn.model_selection import train_test_split from
sklearn.metrics import mean_squared_error, r2_score

%matplotlib inline
```

1.2 Working on Train dataframe

```
[2]: traindf = pd.read_csv('train.csv')
```

```
[3]: traindf.columns
```

```
[3]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea',
'Street',
'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
'BldgType',
'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt',
'YearRemodAdd',
'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd',
'MasVnrType',
'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation',
'BsmtQual',
'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
'Heating',
'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF',
'2ndFlrSF',
'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath',
'FullBath',
'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
```

```

'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu',
'GarageType',
'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea',
'GarageQual',
'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
'PoolQC',
'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
'SaleCondition', 'SalePrice'],
dtype='object')

```

```

[4]: numeric_df = traindf.select_dtypes(include='number')
correlation_matrix = numeric_df.corr()
correlation_matrix['SalePrice'].sort_values(ascending =
False)

```

```

[4]: SalePrice      1.000000
OverallQual      0.790982
GrLivArea        0.708624
GarageCars        0.640409
GarageArea        0.623431
TotalBsmtSF       0.613581
1stFlrSF          0.605852
FullBath          0.560664
TotRmsAbvGrd     0.533723
YearBuilt         0.522897
YearRemodAdd      0.507101
GarageYrBlt       0.486362
MasVnrArea        0.477493
Fireplaces        0.466929
BsmtFinSF1        0.386420
LotFrontage       0.351799
WoodDeckSF        0.324413
2ndFlrSF          0.319334
OpenPorchSF       0.315856
HalfBath          0.284108
LotArea           0.263843
BsmtFullBath      0.227122
BsmtUnfSF         0.214479
BedroomAbvGr      0.168213
ScreenPorch       0.111447
PoolArea          0.092404
MoSold            0.046432
3SsnPorch         0.044584
BsmtFinSF2        -
                  0.011378
BsmtHalfBath      -
                  0.016844

```

```

MiscVal      -
              0.021190
Id           -
              0.021917
LowQualFinSF -
              0.025606
YrSold       -
              0.028923
OverallCond  -
              0.077856
MSSubClass   -
              0.084284
EnclosedPorch -
              0.128578
KitchenAbvGr -
              0.135907
Name: SalePrice, dtype: float64

```

```

[5]: req_tr = _
      ["GarageArea", "OverallQual", "TotalBsmtSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF", "GrLivArea",

```

```

[6]: selected_tr = traindf[req_tr]

```

```

[7]: selected_tr.loc[:, 'TotalBath'] =
      (selected_tr['BsmtFullBath'].fillna(0)
      +
      selected_tr['BsmtHalfBath'].fillna(0)
      + selected_tr['FullBath'].fillna(0) +
      selected_tr['HalfBath'].fillna(0))

      selected_tr.loc[:, 'TotalSF'] = (selected_tr['TotalBsmtSF'].fillna(0)
      +
      selected_tr['1stFlrSF'].fillna(0) +
      selected_tr['2ndFlrSF'].fillna(0) +
      selected_tr['LowQualFinSF'].fillna(0) +
      selected_tr['GrLivArea'].fillna(0))

```

C:\Users\jagan\AppData\Local\Temp\ipykernel_4956\2208519363.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a
 DataFrame. Try using `.loc[row_indexer,col_indexer] = value`
 instead

See the caveats in the documentation:

[https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#r](https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
 eturning-a-view-versus-a-copy `selected_tr.loc[:, 'TotalBath'] =`

```
(selected_tr['BsmtFullBath'].fillna(0) +
C:\Users\jagan\AppData\Local\Temp\ipykernel_4956\2208519363.py:6:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a
DataFrame. Try using .loc[row_indexer,col_indexer] = value
instead
```

See the caveats in the documentation:
[https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#r](https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
returning-a-view-versus-a-copy selected_tr.loc[:, 'TotalSF'] =
(selected_tr['TotalBsmtSF'].fillna(0) +

```
[8]: selected_tr
```

```
[8]:
```

	GarageArea	OverallQual	TotalBsmtSF	1stFlrSF	2ndFlrSF	LowQualFinSF	\
0	548	7	856	856	854	0	
1	460	6	1262	1262	0	0	
2	608	7	920	920	866	0	
3	642	7	756	961	756	0	
4	836	8	1145	1145	1053	0	
...	
1455	460	6	953	953	694	0	
1456	500	6	1542	2073	0	0	
1457	252	7	1152	1188	1152	0	
1458	240	5	1078	1078	0	0	
1459	276	5	1256	1256	0	0	
	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	TotRmsAbvGrd	\
0	1710	1	0	2	1	8	
1	1262	0	1	2	0	6	
2	1786	1	0	2	1	6	
3	1717	1	0	1	0	7	
4	2198	1	0	2	1	9	
...	
1455	1647	0	0	2	1	7	
1456	2073	1	0	2	0	7	
1457	2340	0	0	2	0	9	
1458	1078	1	0	1	0	5	
1459	1256	1	0	1	1	6	
	SalePrice	TotalBath	TotalSF				

0	208500	4	4276
1	181500	3	3786
2	223500	4	4492
3	140000	2	4190
4	250000	4	5541
...
1455	175000	3	4247
1456	210000	3	5688
1457	266500	2	5832
1458	142125	2	3234
1459	147500	3	3768

[1460 rows x 15 columns]

1.3 Keeping only the necessary columns

```
[9]: train_df = _
      ↪selected_tr[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual', 'SalePrice']]
```

```
[10]: train_df
```

```
[10]: TotRmsAbvGrd TotalBath GarageArea TotalSF OverallQual SalePrice
0          8          4          548      4276          7      208500
1          6          3          460      3786          6      181500
2          6          4          608      4492          7      223500
3          7          2          642      4190          7      140000
4          9          4          836      5541          8      250000
...
1455        7          3          460      4247          6      175000
1456        7          3          500      5688          6      210000
1457        9          2          252      5832          7      266500
1458        5          2          240      3234          5      142125
1459        6          3          276      3768          5      147500
```

[1460 rows x 6 columns]

1.4 Splitting the dataset and Creating Pipeline

```
[11]: from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(train_df, test_size =
0.2, random_state = 42) print(f"Rows in train set:
{len(train_set)}\nRows in test set: {len(test_set)}\n")
```

```
Rows in train set: 1168
Rows in test set: 292
```

```
[12]: housing = train_set.drop("SalePrice", axis=1)
housing_labels = train_set["SalePrice"].copy()
```

```
[13]: from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
my_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler())
])
```

```
[14]: X_train = my_pipeline.fit_transform(housing)
```

```
[15]: X_train
```

```
[15]: array([[ -0.96456591, -0.48377079, -0.86383727, -0.13352109, -
 0.82044456], [ 0.27075534, 0.61127627, -0.45626397, -0.13428593, -
 0.08893368],
  [-1.58222654, -1.57881784, -2.25716927, -1.32207838, -
 0.82044456],
 ...,
  [-0.96456591, -0.48377079, 0.45366713, -1.16605156, -
 0.82044456],
  [ 0.27075534, -0.48377079, -1.23349678, -0.26966215,
 0.64257719],
  [ 0.27075534, -0.48377079, 0.87071888, 0.28025593,
 0.64257719]])
```

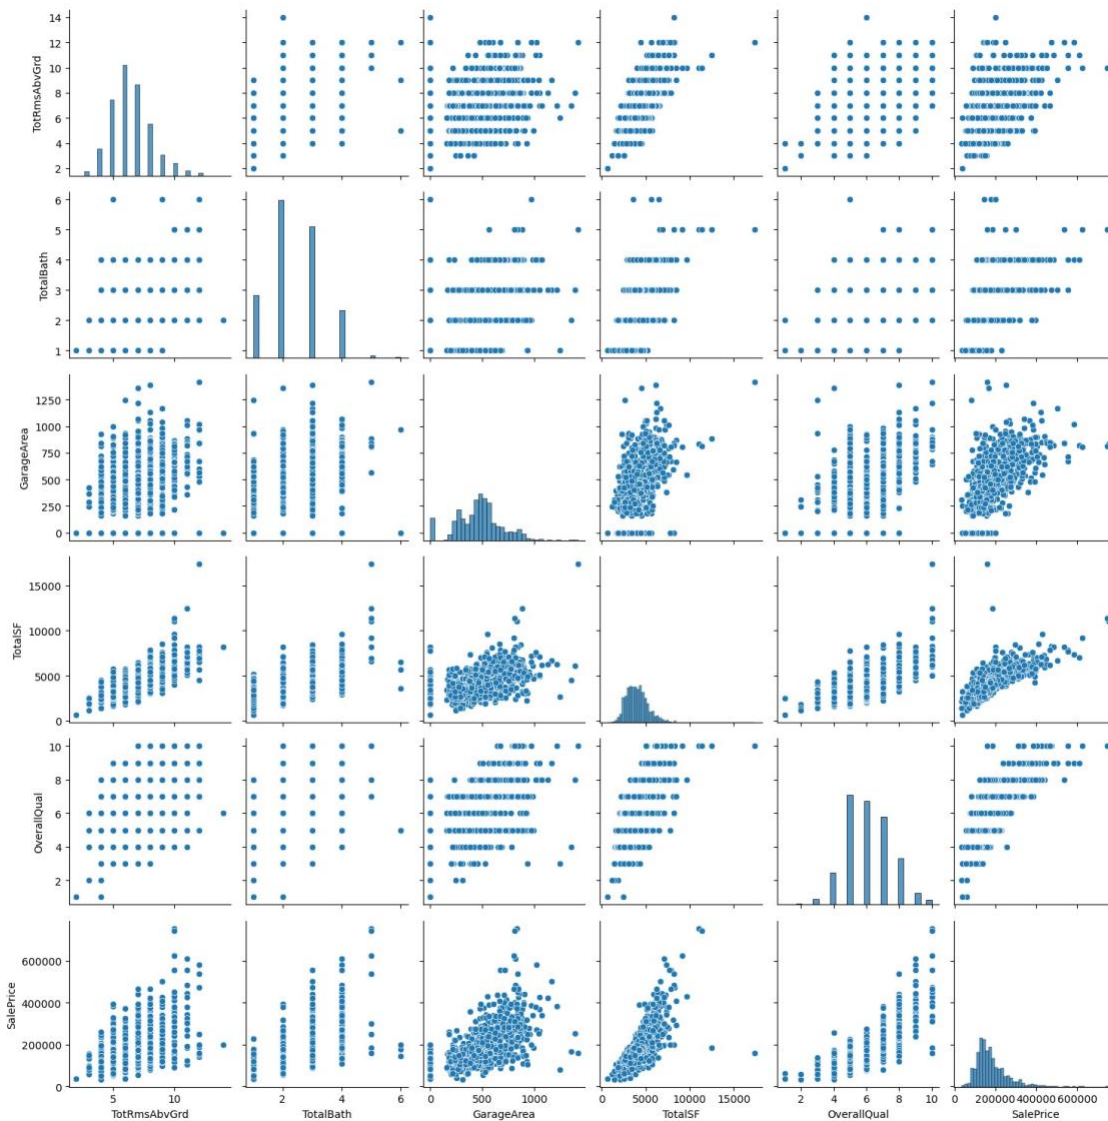
```
[16]: Y_train = housing_labels
```

```
[17]: Y_train.shape
```

```
[17]: (1168,)
```

1.5 Correlations

```
[18]: import warnings
warnings.filterwarnings("ignore", category=UserWarning)
%matplotlib inline
sns.pairplot(train_df)
plt.tight_layout()
plt.show()
```



```
[19]: corr_matrix = train_df.corr()
corr_matrix['SalePrice'].sort_values(ascending = False)
```

```
[19]: SalePrice    1.000000
OverallQual    0.790982
TotalSF        0.773909
```


```

GarageArea      0.623431
TotalBath       0.613005
TotRmsAbvGrd    0.533723
Name: SalePrice, dtype: float64

```

```
[20]: sns.heatmap(train_df.corr(),annot = True)
```

```
[20]: <Axes: >
```



1.6 Working with Test Dataframe

```
[21]: testdf = pd.read_csv("test.csv")
```

```
[22]: testdf.head()
```

```

[22]:  Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0   1461      20    RH      80.0  11622 Pave  NaN    Reg
1   1462      20    RL      81.0  14267 Pave  NaN    IR1
2   1463      60    RL      74.0  13830 Pave  NaN    IR1
3   1464      60    RL      78.0   9978 Pave  NaN    IR1

```



```

4  1465      120   RL    43.0  5005  Pave  NaN   IR1

LandContour Utilities ... ScreenPorch PoolArea PoolQC Fence
MiscFeature \
0    Lvl  AllPub ...   120   0    NaN MnPrv  NaN  1 Lvl  AllPub ...
0    0    NaN   NaN   Gar2
2    Lvl  AllPub ...    0    0    NaN MnPrv  NaN  3 Lvl  AllPub ...
0    0    NaN   NaN   NaN
4      HLS  AllPub ...      144      0    NaN   NaN      NaN

MiscVal MoSold YrSold SaleType SaleCondition
0      0    6      2010 WD      Normal
1     12500    6      2010 WD      Normal
2      0    3      2010 WD      Normal
3      0    6      2010 WD      Normal
4      0    1      2010 WD      Normal

```

[5 rows x 80 columns]

```

[23]: req_tst = _
      4["GarageArea", "OverallQual", "TotalBsmtSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF", "GrLivArea",

```

```

[24]: selected_tst = testdf[req_tst]

```

```

[25]: selected_tst.loc[:, 'TotalBath'] =
(selected_tst['BsmtFullBath'].fillna(0) +
selected_tst['BsmtHalfBath'].fillna(0)
+ selected_tst['FullBath'].fillna(0) +
selected_tst['HalfBath'].fillna(0))

selected_tst.loc[:, 'TotalSF'] =
(selected_tst['TotalBsmtSF'].fillna(0) +
selected_tst['1stFlrSF'].fillna(0) +
selected_tst['2ndFlrSF'].fillna(0) +
selected_tst['LowQualFinSF'].fillna(0)
+ selected_tst['GrLivArea'].fillna(0))

```

C:\Users\jagan\AppData\Local\Temp\ipykernel_4956\771691818.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame. Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation:

https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy selected_tst.loc[:, 'TotalBath'] =

```
(selected_tst['BsmtFullBath'].fillna(0) +
C:\Users\jagan\AppData\Local\Temp\ipykernel_4956\771691818.py:6:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a
DataFrame. Try using .loc[row_indexer,col_indexer] = value
instead
```

See the caveats in the documentation:

[https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#r](https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
[eturning-a-view-versus-a-copy](https://pandas.pydata.org/pandasdocs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
selected_tst.loc[:, 'TotalSF'] = (selected_tst['TotalBsmtSF'].fillna(0) +
[26]: selected_tst
```

```
[26]:  GarageArea  OverallQual  TotalBsmtSF  1stFlrSF  2ndFlrSF  LowQualFinSF  \
0          730.0    5      882.0 896    0    0
1          312.0    6     1329.0    1329  0    0
2          482.0    5     928.0 928    701  0
3          470.0    6     926.0 926    678  0
4          506.0    8     1280.0    1280  0    0
...          ...          ...          ...          ...          ...          ...
1454         0.0    4     546.0 546    546  0
1455        286.04     546.0 546    546  0
1456        576.05     1224.0    1224  0    0
1457         0.0    5     912.0 970    0    0
1458        650.07     996.0 996    1004  0

      GrLivArea  BsmtFullBath  BsmtHalfBath  FullBath  HalfBath  TotRmsAbvGrd  \
0           896      0.0    0.0    1    0    5
1          1329      0.0    0.0    1    1    6
2          1629      0.0    0.0    2    1    6
3          1604      0.0    0.0    2    1    7
4          1280      0.0    0.0    2    0    5
...          ...          ...          ...          ...          ...
1454        1092      0.0    0.0    1    1    5
1455        1092      0.0    0.0    1    1    6
1456        1224      1.0    0.0    1    0    7
1457         970      0.0    1.0    1    0    6
1458        2000      0.0    0.0    2    1    9
```

```
TotalBath  TotalSF
```

```

0      1.0   2674.0 1
2.0   3987.0 2    3.0
4186.0 3    3.0
4134.0
4      2.0   3840.0
...
1454      2.0   2730.0
1455      2.0   2730.0
1456      2.0   3672.0
1457      2.0   2852.0
1458      3.0   4996.0
[1459 rows x 14 columns]

```

```
[27]: test_df_unproc = 
↳selected_tst[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']]
```

```
[28]: test_df_unproc
```

```
[28]:   TotRmsAbvGrd TotalBath GarageArea TotalSF OverallQual
0           5 1.0    730.0 2674.0      5
1           6 2.0    312.0 3987.0      6
2           6 3.0    482.0 4186.0      5
3           7 3.0    470.0 4134.0      6
4           5 2.0    506.0 3840.0      8
...
1454          5 2.0     0.0  2730.0      4
1455          6 2.0    286.0 2730.0      4
1456          7 2.0    576.0 3672.0      5
1457          6 2.0     0.0  2852.0      5
1458          9 3.0    650.0 4996.0      7
[1459 rows x 5 columns]
```

```
[29]: test_df = test_df_unproc.fillna(test_df_unproc.mean())
```

```
[30]: x_test = my_pipeline.
↳transform(test_df[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']]).
↳values)
```

```
[31]: x_test
```

```
[31]: array([[ -0.96456591, -1.57881784, 1.2024646 , -1.10333489, -
0.82044456], [-0.34690528, -0.48377079, -0.77853123, -0.09910341, -
0.08893368],
[-0.34690528, 0.61127627, 0.02713693, 0.05309923, -0.82044456],
...,
[ 0.27075534, -0.48377079, 0.47262403, -0.34002719, -
0.82044456],
[-0.34690528, -0.48377079, -2.25716927, -0.96719384, -
0.82044456],
[ 1.50607659, 0.61127627, 0.82332664, 0.67261751, 0.64257719]])
```

1.7 Model Selection

```
[32]: #model = LinearRegression()
#model = DecisionTreeRegressor()
model = RandomForestRegressor()
model.fit(X_train,Y_train)
```

```
[32]: RandomForestRegressor()
```

```
[33]: y_train_pred = model.predict(X_train)
```

```
[34]: y_train_pred[:5]
```

```
[34]: array([145283. , 172053.95, 89698. , 168814. , 141625. ])
```

```
[35]: some_data = housing.iloc[:5]
some_labels = housing_labels.iloc[:5]
```

```
[36]: proc_data = my_pipeline.transform(some_data)
```

```
[37]: model.predict(proc_data)
```

```
[37]: array([145283. , 172053.95, 89698. , 168814. , 141625. ])
```

```
[38]: list(some_labels)
```

```
[38]: [145000, 178000, 85000, 175000, 127000]
```

```
[39]: train_mse = mean_squared_error(Y_train,y_train_pred)
```

```
[40]: train_rmse = np.sqrt(train_mse)
```

```
[41]: print(f"Training MSE: {train_mse:.2f}, Training RMSE: {train_rmse:.2f}")
```

Training MSE: 182754306.10, Training RMSE: 13518.67

1.8 Cross - Validation

```
[42]: from sklearn.model_selection import cross_val_score
scores = _
    ↪ cross_val_score(model,X_train,Y_train,scoring="neg_mean_squared_error",cv = _
    ↪ 200)
rmse_scores = np.sqrt(-scores)
```

```
[43]: rmse_scores

[43]: array([ 23277.76446447, 13031.67234441, 24491.16208088,
12916.01052409, 46110.0745834 , 12685.13165022, 18886.3441575 ,
11071.47547444,
11616.79173023, 52066.94494877, 31510.80666534,
29639.23399926,
12151.07107851, 9319.99965266, 17079.07143005, 20884.17467223,
18824.49033057, 33193.65937336, 36854.67178624,
23093.95361294,
23292.9684995 , 19432.72111562, 18547.99982299,
27613.27744408,
20402.00284752, 17822.98485906, 42015.72882012,
36937.17830823,
175165.06203852, 49690.4027053 , 23330.91946822, 31513.45824178,
22887.73799821, 31010.13490391, 47989.62053758,
13329.38499857,
21928.82341947, 28502.91443722, 17245.77254429,
31486.07563826,
22230.82920158, 32465.60889805, 27843.04913432,
35600.14766459, 33868.94917467, 31457.61022018, 27956.1485217
, 40078.26812993,
20590.99533685, 21056.55475528, 19898.88278442,
50438.13615389,
39470.50551592, 35012.49792537, 22559.83629881,
25338.56615794,
10921.06969935, 27397.02810506, 25766.11480393,
31086.04232992,
201242.98515478, 12028.23570696, 26659.38149476, 39101.03975416,
22930.10835334, 25187.16704816, 41651.52119048,
29282.34773083,
8224.18203629, 14246.79672003, 56886.03402084, 27417.3801067
, 55466.02380919, 18062.31089481, 46371.11498073,
45814.2412921 ,
37923.06090531, 29540.32205912, 40406.01041065,
33253.30878663,
28357.63588478, 31339.75470533, 106127.87363791,
10572.52219235,
```

44907.81340181, 31278.70874266, 17627.15366185,
 23986.51913983,
 34952.10095183, 82162.55603195, 20900.97615594, 23172.9641621
 ,
 26814.03423907, 24597.32265251, 20439.71940229,
 22240.29752072,
 23770.5040223 , 28314.47719879, 69329.96980133,
 13772.33040116,
 32481.10212308, 29130.99328727, 44167.46596499,
 53298.49040065,
 18883.32266379, 17803.70762922, 35231.98484315,
 20413.09211208,
 18867.41154384, 14814.45213528, 19582.28398618,
 11716.30128734,
 10636.0999438 , 22564.76237833, 31336.81397651,
 31644.56494109,
 75496.23776138, 30016.20028031, 20454.74205292,
 30066.09581106,
 30854.08791016, 25820.47922475, 13727.55889958,
 19292.76478605,
 30038.27850124, 20108.84141508, 27973.08744932,
 45101.61936736,
 38876.15567145, 16128.61505828, 33935.44165938, 8855.15987218,
 25776.08507509, 25150.93116638, 33212.31806431,
 14352.15252709,
 44717.67500818, 32649.13731135, 28537.97918331,
 12465.11203373,
 26243.67447064, 25454.07721303, 27022.5697848 ,
 10317.17889282,
 21891.22473637, 29261.49148125, 33452.4737976 , 13743.8454676
 ,
 11277.83822064, 16237.91736661, 24039.63841877, 26656.8558185
 ,
 15413.36324152, 42944.44827572, 31109.2988122 ,
 15811.60360654,
 25686.08533944, 25779.34969772, 20482.6668366 ,
 25274.52807206,
 53481.51708005, 27726.34633021, 27098.11067491,
 18400.37491698,
 51419.98502293, 20869.90451782, 12858.92838191,
 61556.68668169,
 21786.96406963, 31011.97109104, 33354.55797674,
 22402.05803861,
 14163.05024854, 14270.74072466, 19769.06816004, 17101.5888594
 ,

```

24282.57518006, 17461.77872288, 38024.22855112,
20424.51415501,
36820.8916111 , 9820.05229768, 15945.06899661, 6109.83019202,
19796.08373925, 29891.36122746, 26424.62050824,
11914.26957176,
27561.72469719, 42619.58945645, 27378.43479366,
21589.13644804,
17407.48298693, 13255.09382259, 16332.29125757,
13944.69038795,
62926.75609394, 19634.57473268, 21195.25876501,
32191.88871351])

```

```

[44]: def print_scores(scores):
      print("Scores:", scores)
      print("Mean:", scores.mean())
      print("Standard Deviation", scores.std())

```

```

[45]: print_scores(rmse_scores)

```

```

Scores: [ 23277.76446447 13031.67234441 24491.16208088 12916.01052409
46110.0745834 12685.13165022 18886.3441575 11071.47547444
11616.79173023 52066.94494877 31510.80666534 29639.23399926
12151.07107851 9319.99965266 17079.07143005 20884.17467223
18824.49033057 33193.65937336 36854.67178624 23093.95361294
23292.9684995 19432.72111562 18547.99982299 27613.27744408
20402.00284752 17822.98485906 42015.72882012 36937.17830823
175165.06203852 49690.4027053 23330.91946822 31513.45824178
22887.73799821 31010.13490391 47989.62053758 13329.38499857
21928.82341947 28502.91443722 17245.77254429 31486.07563826
22230.82920158 32465.60889805 27843.04913432 35600.14766459
33868.94917467 31457.61022018 27956.148521740078.26812993
20590.99533685 21056.55475528 19898.88278442 50438.13615389
39470.50551592 35012.49792537 22559.83629881 25338.56615794
10921.06969935 27397.02810506 25766.11480393 31086.04232992
201242.98515478 12028.23570696 26659.38149476 39101.03975416
22930.10835334      25187.16704816      41651.52119048
29282.34773083      8224.18203629      14246.79672003
56886.03402084      27417.3801067      55466.02380919
18062.31089481 46371.11498073 45814.2412921
37923.06090531 29540.32205912 40406.01041065 33253.30878663
28357.63588478 31339.75470533 106127.87363791 10572.52219235
44907.81340181 31278.70874266 17627.15366185 23986.51913983
34952.10095183 82162.55603195 20900.97615594 23172.9641621
26814.03423907 24597.32265251 20439.71940229 22240.29752072
23770.5040223 28314.47719879 69329.96980133 13772.33040116
32481.10212308 29130.99328727 44167.46596499 53298.49040065
18883.32266379 17803.70762922 35231.98484315 20413.09211208

```

```

18867.41154384 14814.45213528 19582.28398618 11716.30128734
10636.0999438 22564.76237833 31336.81397651 31644.56494109
75496.23776138 30016.20028031 20454.74205292 30066.09581106
30854.08791016 25820.47922475 13727.55889958 19292.76478605
30038.27850124 20108.84141508 27973.08744932 45101.61936736
38876.15567145 16128.61505828 33935.441659388855.15987218
25776.08507509 25150.93116638 33212.31806431 14352.15252709
44717.67500818 32649.13731135 28537.97918331 12465.11203373
26243.67447064 25454.07721303 27022.569784810317.17889282
21891.22473637 29261.49148125 33452.473797613743.8454676
11277.83822064 16237.91736661 24039.63841877 26656.8558185
15413.36324152 42944.44827572 31109.298812215811.60360654
25686.08533944 25779.34969772 20482.666836625274.52807206
53481.51708005 27726.34633021 27098.11067491 18400.37491698
51419.98502293 20869.90451782 12858.92838191 61556.68668169
21786.96406963 31011.97109104 33354.55797674 22402.05803861
14163.05024854 14270.74072466 19769.06816004 17101.5888594
24282.57518006 17461.77872288 38024.22855112 20424.51415501
36820.8916111 9820.05229768 15945.06899661 6109.83019202
19796.08373925 29891.36122746 26424.62050824 11914.26957176
27561.72469719 42619.58945645 27378.43479366 21589.13644804
17407.48298693 13255.09382259 16332.29125757 13944.69038795
62926.75609394 19634.57473268 21195.25876501 32191.88871351]
Mean: 29103.74552461538
Standard Deviation 21139.34340094215

```

```
[46]: y_pred=model.predict(x_test)
```

```
[47]: y_pred
```

```
[47]: array([129596.33, 155331.14, 148147.5 , ..., 138024.85, 107234.5 ,
          231268. ])
```

```
[48]: pred=pd.DataFrame(y_pred)
sub_df=pd.read_csv('sample_submission.csv')
datasets=pd.concat([sub_df['Id'],pred],axis=1)
datasets.columns=['Id','SalePrice']
datasets.to_csv('sample_submission.csv',index=False)
```